
Internship Report

Karachi AQI Predictor

Air Quality Index Prediction System
Using Machine Learning & MLOps

Submitted by: Arham Altaf

Institution: IBA Karachi

Host Organisation: 10 Pearls

Report Date: June 2026

Contents

1	Introduction & Motivation	2
1.1	The Dataset Decision	2
2	Dataset	3
2.1	Overview	3
2.2	Feature Groups	3
2.3	AQI Category Distribution	3
3	System Architecture	3
3.1	Technology Stack	4
4	Machine Learning Models	4
4.1	Validation Strategy	4
4.2	Model Performance	5
4.3	Performance Visualisation	6
4.4	Model Registry	6
5	Exploratory Data Analysis (EDA)	6
5.1	Feature Engineering Pipeline	6
6	Automation & MLOps	7
6.1	GitHub Actions Workflows	7
6.2	Why Automate?	7
7	Challenges & Problems Faced	7
8	Key Learnings	8
9	Project Structure	8
10	SHAP Analysis & Model Interpretability	9
10.1	Core Concept	9
10.2	Why SHAP Matters for This Project	9
10.3	Key Findings	10
10.4	Individual Prediction Explanation (Waterfall Chart)	10
10.5	Technical Implementation	10
11	Conclusion	10

1 Introduction & Motivation

This project, designed by 10Pearls, was fundamentally different from the typical Kaggle-based assignments I had completed during my Machine Learning course at IBA Karachi. Instead of working with a preprocessed and ready-made dataset, this project required building a complete end-to-end pipeline, starting from real-time data collection through APIs to feature engineering, model training, and deployment. This made the project significantly more challenging and closer to real-world machine learning systems.

The core prediction task is to forecast the Air Quality Index (AQI) category for Karachi, Pakistan up to 72 hours in advance. Understanding air quality is a meaningful challenge because AQI is influenced by two fundamentally different types of factors:

1. **Pollutants:** particulate matter such as PM2.5, which directly determines air quality ratings.
2. **Weather conditions:** temperature, humidity, wind speed, precipitation, and atmospheric pressure, which influence how pollutants disperse or accumulate.

The realization that both of these dimensions had to be captured, combined, and modelled together made this project genuinely educational beyond the classroom.

1.1 The Dataset Decision

An important design evolution in this project was the definition of the target variable for AQI prediction. Two approaches were evaluated:

- **Method A (Categorical API):** Using the OpenWeather API's pre-calculated AQI category (values 1–5) directly as the prediction target for a classification model.
- **Method B (PM2.5 Regression):** Predicting the raw, continuous PM2.5 concentration value using regression models, subsequently mapping this value to a standardized AQI scale (0–500) via the official US EPA formula.

Initially, Method A was favored for its simplicity. By utilizing the API's pre-processed index, the pipeline remained straightforward, focusing on feature engineering without the overhead of additional mathematical transformations. This allowed for rapid prototyping and efficient computational performance, providing a practical baseline for forecasting. However, as the project matured, the limitations of Method A became apparent. Predicting an abstracted 1–5 index led to a significant loss of variance, as distinct physical pollutant levels were compressed into the same category, limiting the model's ability to learn nuanced environmental patterns. Consequently, the architecture was transitioned to Method B. This approach is significantly more rigorous, treating the problem as a regression task to predict raw PM2.5 concentrations ($\mu g/m^3$). While Method B introduced higher complexity, requiring a transition from classification to regression models, the implementation of the EPA conversion formula, and the use of regression specific metrics like RMSE and MAE, it offers vastly improved reliability and interpretability. By training on continuous physical data, the model can capture high-resolution atmospheric trends that

are otherwise masked by categorical thresholds. This evolution from Method A to Method B ensures the system adheres to industry standard environmental forecasting practices, resulting in a more robust and physically meaningful prediction engine.

2 Dataset

2.1 Overview

Table 1: Dataset Characteristics

Property	Value
Location	Karachi, Pakistan (24.8608°N, 67.0104°E)
Duration	78 days
Frequency	Hourly
Total Records	1,872 hourly data points
Feature Count	30 engineered features
Target Variable	PM 2.5
Data Sources	Open-Meteo (weather) + Open Weather (AQI)

2.2 Feature Groups

The 30 features were carefully engineered across several categories to capture both instantaneous conditions and historical patterns.

Table 2: Engineered Feature Categories

Category	Features	Description
Raw Weather	5	Temperature, humidity, wind speed, wind direction, cloud cover
Raw Pollution	8	PM2.5, PM10, CO, NO, NO ₂ , O ₃ , SO ₂ , NH ₃ (all in $\mu\text{g}/\text{m}^3$)
Target	1	AQI category (from Open Weather API)
Time-Based	5	Hour, day, month, day of week, cyclical encodings (sin / cos)
Lag Features	3	AQI values at 24h, 48h, 72h ago (prevents data leakage)
Rolling Statistics	6+	72-hour window means and standard deviations
Derived Features	8+	Weather interactions, trend indicators

2.3 AQI Category Distribution

The target variable has four classes, each corresponding to an air quality level:

3 System Architecture

The system is designed as a fully automated MLOps pipeline with no manual intervention required after initial setup.




Figure 1: Approximate distribution of AQI categories across 1,872 hourly records

Table 3: Technology Stack

Component	Technology	Purpose
Data Collection	Open-Meteo API, OpenWeather API	Live weather & AQI data
Feature Engineering	Python (pandas, numpy)	Lag, rolling, derived features
Database	MongoDB Atlas	Feature & model storage
ML Models	scikit-learn, XGBoost, Light-GBM	Classification
Automation	GitHub Actions	Hourly & daily workflows
Serving	Streamlit Cloud	Interactive dashboard
Model Persistence	Joblib (.joblib files)	Local model serialization

3.1 Technology Stack

4 Machine Learning Models

4.1 Validation Strategy

To ensure a robust evaluation of model performance across diverse atmospheric conditions, the dataset was split using a stratified random shuffle. Unlike a simple temporal split, this approach guarantees that each AQI class is represented proportionately in both the training and testing sets, which is critical for mitigating issues related to class imbalance in air quality datasets.

Table 4: Train/Test Split

Split	Records	Proportion
Training Set	1,440	80%
Test Set	360	20%
Cross-Validation	5-fold stratified	On training set




Figure 2: End-to-end MLOps pipeline architecture

4.2 Model Performance

Three regression models were trained and evaluated using Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), and Coefficient of Determination (R^2). XGBoost achieved the best overall performance, obtaining the lowest prediction errors and the highest R^2 score.

Table 5: Regression Model Performance Comparison

Model	RMSE	MAE	R^2 Score
XGBoost	1.34	0.90	0.97
LightGBM	1.47	0.95	0.96
Random Forest	3.04	2.25	0.85




Figure 3: Comparison of RMSE, MAE, and R^2 scores across regression models

4.3 Performance Visualisation

4.4 Model Registry

All three models are tracked in a MongoDB model registry with full metadata including hyperparameters, performance metrics, and training timestamps. They are also serialized locally as joblib files for fast inference.

5 Exploratory Data Analysis (EDA)

Before modelling, extensive EDA was performed to understand the data and inform feature engineering decisions.

5.1 Feature Engineering Pipeline

The EDA phase revealed that raw hourly readings alone were insufficient. Key insights:

- **AQI conditions are autocorrelated:** what the air quality is now is strongly related to what it was 24, 48, and 72 hours ago.
- **Cyclical time encodings** (sin / cos of hour and day) better capture daily and weekly periodicity than raw integer values.
- **Rolling statistics** over a 72-hour window capture the trend and variability of conditions, not just the snapshot value.
- **Weather interactions** (e.g., high humidity + low wind $\rightarrow$ poor dispersion) can be captured as derived features.




Figure 4: Illustrative diurnal AQI pattern showing typically better air quality midday (higher dispersion) and worse quality at night and early morning

6 Automation & MLOps

6.1 GitHub Actions Workflows

A key learning was how to use GitHub Actions to automate repetitive tasks—essentially the ML equivalent of a `for` loop that runs on a schedule in the cloud.

Table 6: GitHub Actions Workflows

Workflow	Schedule	Actions
hourly-data-collection.yml	Every hour at :00	Fetch weather + AQI from APIs, engineer features, store in MongoDB
daily-model-training.yml	Daily at 2:00 AM UTC	Train all 3 models, evaluate, update model registry, commit joblib files to Git

6.2 Why Automate?

- **Data freshness:** Hourly collection ensures the model always has recent data.
- **Model freshness:** Daily retraining prevents model drift as Karachi’s seasonal pollution patterns evolve.
- **Zero manual effort:** After initial setup, the system runs indefinitely without human intervention.

7 Challenges & Problems Faced

1. **Building a Dataset from Scratch:** Unlike the Kaggle competition in class where the dataset was provided, this project required sourcing, cleaning, and combining data from two entirely separate APIs (Open-Meteo for weather and OpenWeather for AQI). Ensuring

the timestamps aligned correctly between the two sources required careful engineering. The realization of how much effort goes into real-world data collection was a major learning moment.

2. **Choosing the Right AQI Prediction Method:** Deciding between predicting AQI as a direct categorical output vs. first predicting PM2.5 and then converting it using the EPA formula was confusing. The PM2.5 route introduced compounded error: regression errors in PM2.5 would propagate into the AQI classification. Ultimately, using the API's categorical AQI output directly was simpler and more reliable, but arriving at this decision required understanding how AQI is actually computed.
3. **Understanding GitHub Actions:** Scheduling automated workflows was a completely new concept. Moving from "run this Python script manually" to "this script runs every hour in the cloud automatically" required understanding YAML workflow syntax, GitHub Secrets for API keys, and how to trigger jobs on a cron schedule. The first workflows failed several times before they ran successfully end-to-end.
4. **MongoDB Integration:** Connecting the feature pipeline, the model trainer, and the Streamlit dashboard to the same MongoDB Atlas database without conflicts—especially with concurrent reads and writes from GitHub Actions and the dashboard—required careful database schema design and proper use of connection strings and Streamlit secrets management.

8 Key Learnings

1. **Real-world data is hard:** Creating a dataset is far more work than using a pre-made one. API rate limits, missing values, timestamp mismatches, and schema design are all real engineering problems.
2. **Feature engineering matters more than model choice:** The lag and rolling features contributed more to model performance than switching between RF, XGBoost, and LightGBM.
3. **Time-series requires special care:** Temporal splitting and lag features are non-negotiable for honest evaluation.
4. **MLOps closes the loop:** Training a good model is only half the job. Automated pipelines, model registries, and deployment are what make ML useful in the real world.
5. **GitHub Actions = scheduled automation:** A cron job in the cloud that keeps your data and models fresh without any manual work.

9 Project Structure

```
karachi-aqi-predictor/  
  .github/workflows/
```

```

    hourly-data-collection.yml    # Hourly feature collection
    daily-model-training.yml      # Daily retraining
src/
    config.py                    # Configuration management
    database.py                  # MongoDB handler
    feature_pipeline.py          # Hourly data collection & engineering
    training_pipeline.py         # Model training & evaluation
    model_registry.py           # Model loading & predictions
models/                          # Serialized models
    randomforest_model.joblib
    xgboost_model.joblib
    lightgbm_model.joblib
    scaler.joblib
    label_encoder.joblib
    feature_columns.joblib
app.py                          # Streamlit dashboard
requirements.txt
README.md

```

10 SHAP Analysis & Model Interpretability

Traditional machine learning models like LightGBM, XGBoost, and Random Forest are often called "black boxes"—they make accurate predictions but offer no explanation of why. SHAP (SHapley Additive exPlanations) solves this by computing exactly how much each feature contributes to each individual prediction.

10.1 Core Concept

SHAP is grounded in Shapley values from cooperative game theory. Each feature is treated as a "player" in a game, and its contribution is calculated by:

1. Computing predictions with and without that feature across all possible feature combinations.
2. Averaging contributions fairly so no feature is over- or under-credited.
3. Guaranteeing that all SHAP values sum exactly to the difference between the model's prediction and the baseline (average AQI).

10.2 Why SHAP Matters for This Project

Table 7: Benefits of SHAP Interpretability

Benefit	What it Provides
Transparency & Trust	Understand which environmental factors actually drive AQI predictions. Validate that the model uses sensible patterns (e.g., high PM2.5 $\rightarrow$ worse AQI).
Debugging	Detect if the model relies on spurious or accidental correlations.
Scientific Insight	Discover which factors matter most for Karachi's air quality.

10.3 Key Findings

What SHAP Revealed About the Model

1. **Weather dominates:** `cloud_cover` (SHAP: 2.45) is the single most influential feature, followed by temperature and `wind_direction`. Weather conditions control how pollutants disperse in the atmosphere.
2. **Temporal patterns matter:** `day_of_week_sin` (SHAP: 1.98) and `hour_sin` capture human behavioural patterns—traffic, industry, and cooking cycles—that vary by time of day and day of week.
3. **Lag features are critical:** `aqi_lag_48h` (1.35), `aqi_lag_72h` (1.23), and `aqi_lag_24h` (1.12) confirm that recent AQI history is a strong predictor of future AQI, validating the feature engineering decisions.
4. **Rolling mean captures trends:** `aqi_rolling_mean_72h` (1.73) reflects multi-day pollution accumulation or dispersal trends.
5. **Raw pollutants rank lower than expected:** `pm2_5` (0.97) and `pm10` (0.55) have lower SHAP values not because they are unimportant, but because their information is already captured by the AQI lag and rolling features, which encode historical pollution levels.

10.4 Individual Prediction Explanation (Waterfall Chart)

Beyond global importance, SHAP also explains individual predictions via a waterfall chart. Starting from the baseline average AQI, each feature either pushes the prediction higher (worse air quality) or lower (better air quality).

10.5 Technical Implementation

This interpretability layer transforms the AQI predictor from a "magic algorithm" into an explainable decision support tool that domain experts and stakeholders can understand, question, and trust.

11 Conclusion

The Karachi AQI Predictor is more than a machine learning model—it is a complete, production-grade MLOps system. Starting from zero data, the project progressed through API integration,




Figure 5: Illustrative waterfall chart for a single prediction. Each bar shows how much a feature shifts the prediction above or below the baseline AQI.

Table 8: SHAP Technical Details

Property	Detail
Algorithm	TreeExplainer (optimised for tree-based models)
Backend	Fast C++ implementation for LightGBM / XGBoost
Consistency	SHAP values guaranteed to sum exactly to the prediction
Caching	Explainer cached for 1 hour on Streamlit to avoid recomputation
Scope	Both global (all predictions) and local (individual predictions)

dataset creation, feature engineering, model training, automated scheduling, database management, and live deployment.

The best model, LightGBM, achieves 93.9% test accuracy and 93.1% cross-validation accuracy, providing reliable 72-hour AQI forecasts for one of South Asia’s most populous cities.

More importantly, the experience of building this end-to-end taught lessons that no pre-packaged Kaggle dataset ever could—from the difficulty of sourcing and cleaning real data, to the engineering discipline required to prevent data leakage, to the satisfaction of watching an automated system run in the cloud without human intervention.